

Interview DBA Part 1

What is a database administrator's main role?

A **database administrator (DBA)** is responsible for managing and maintaining databases. Here are the main roles of a DBA in simple terms:

- **Monitoring Databases:** They check if the databases are working properly and fix any issues that arise.
- **Backing Up Data:** DBAs regularly save copies of the data to prevent loss in case of problems.
- **Optimizing Performance:** They make sure the databases run efficiently, which means they help speed up slow queries and improve overall performance.
- **Managing User Access:** DBAs control who can access the database and what they can do, ensuring security.
- **Planning for Growth:** They prepare for future needs by ensuring the database can handle more users and data as the organization grows.

In summary, a DBA keeps databases running smoothly, secure, and ready for future demands.

What is the most challenging task for a database administrator?

One of the most challenging tasks for a database administrator (DBA) is **troubleshooting and resolving performance issues**. Here's why it can be difficult:

- **Identifying the Problem:** Performance issues can arise from various sources, such as slow queries, hardware limitations, or network problems. Pinpointing the exact cause can be complex.

- **Analyzing Query Performance:** DBAs need to analyze how queries are executed and identify which ones are causing delays. This often requires deep knowledge of the database structure and query optimization techniques.
- **Balancing Resources:** They must ensure that the database has enough resources (like memory and processing power) while also managing costs. This involves making decisions about upgrades or changes to the database environment.
- **Maintaining Uptime:** DBAs must resolve issues quickly to minimize downtime, which can affect business operations. This requires quick thinking and effective problem-solving skills.

Overall, troubleshooting performance issues requires a combination of technical skills, analytical thinking, and experience.

What strategies would you use to handle multiple support tickets?

Handling multiple support tickets efficiently is crucial for a database administrator (DBA). Here are some effective strategies:

- **Prioritization:** Assess the urgency and impact of each ticket. Focus on critical issues that affect many users or business operations first.
- **Categorization:** Group similar tickets together. This allows you to address common issues in batches, saving time and effort.
- **Use of Templates:** Create response templates for frequently asked questions or common problems. This speeds up the response time and ensures consistency.
- **Time Management:** Allocate specific time blocks to work on support tickets. This helps maintain focus and prevents distractions.
- **Documentation:** Keep detailed records of issues and solutions. This can help in resolving similar tickets in the future more quickly.

- **Communication:** Keep users informed about the status of their tickets. Clear communication can help manage expectations and reduce follow-up inquiries.

By implementing these strategies, a DBA can effectively manage multiple support tickets while ensuring timely and quality responses.

How can you apply database lifecycle management in a project?

Applying **Database Life Cycle Management** in a project involves following the four stages effectively. Here's how you can do it:

1. Requirements Analysis:

- **Gather Information:** Talk to stakeholders (like users and developers) to understand what data is needed and how it will be used.
- **Define Goals:** Clearly outline what the database should achieve, such as storing customer information or tracking sales.

2. Design and Plan:

- **Create a Database Model:** Use tools like Entity Relationship Diagrams (ERDs) to visualize how data will be organized (tables, relationships, etc.).
- **Plan Resources:** Determine the necessary hardware and software resources, such as server capacity and storage needs.

3. Implementation:

- **Build the Database:** Set up the database according to the design. Create tables, define relationships, and set up security measures.
- **Import Data:** If needed, migrate existing data into the new database.

4. Monitor and Maintain:

- **Regular Checks:** Continuously monitor the database for performance issues, such as slow queries or security vulnerabilities.

- **Update and Optimize:** Apply updates, optimize queries, and perform regular backups to ensure data safety and efficiency.

By following these steps, you can ensure that your database is well-structured, secure, and meets the needs of your project effectively.

What challenges might arise during the implementation stage?

During the **implementation stage** of Database Life Cycle Management, several challenges may arise:

1. Data Migration Issues:

- **Incompatibility:** Existing data may not fit the new database structure, leading to errors during migration.
- **Data Loss:** There is a risk of losing data if not properly backed up before migration.

2. Performance Problems:

- **Slow Queries:** Newly created queries may not be optimized, leading to slow performance.
- **Resource Allocation:** Insufficient server resources (like memory or storage) can hinder database performance.

3. Security Concerns:

- **Access Control:** Misconfiguring user permissions can lead to unauthorized access or data breaches.
- **Vulnerabilities:** New databases may have security gaps that need to be addressed immediately.

4. Technical Challenges:

- **Configuration Errors:** Incorrect settings during setup can lead to functionality issues.

- **Integration Issues:** Difficulty in integrating the database with existing applications or systems.

5. User Training:

- **Lack of Familiarity:** Users may struggle to adapt to the new system if they are not adequately trained.
- **Resistance to Change:** Some users may be resistant to adopting the new database system.

Addressing these challenges requires careful planning, thorough testing, and effective communication with all stakeholders.

How can you optimize queries during implementation?

Optimizing queries during the implementation stage is crucial for ensuring good database performance. Here are some effective strategies:

1. Use Indexes:

- **Create Indexes:** Indexes can speed up data retrieval. Identify columns that are frequently used in search conditions and create indexes on them.
- **Avoid Over-Indexing:** While indexes improve read performance, too many can slow down write operations. Balance is key.

2. Write Efficient SQL Queries:

- **Select Only Necessary Columns:** Instead of using `SELECT *`, specify only the columns you need to reduce data load.
- **Use WHERE Clauses:** Filter results with `WHERE` clauses to limit the number of rows returned.

3. Optimize Joins:

- **Choose the Right Join Type:** Use INNER JOINS when possible, as they are generally faster than OUTER JOINS.

- **Join on Indexed Columns:** Ensure that the columns used for joining tables are indexed.

4. Analyze Query Execution Plans:

- **Use Execution Plans:** Analyze the execution plan of your queries to understand how the database engine processes them. Look for bottlenecks or inefficient operations.
- **Refactor Queries:** Based on the execution plan, refactor queries to improve performance.

5. Limit Result Sets:

- **Use Pagination:** For large result sets, implement pagination to load data in smaller chunks rather than all at once.
- **Use LIMIT:** If you only need a subset of results, use the `LIMIT` clause to restrict the number of rows returned.

6. Regularly Monitor Performance:

- **Track Query Performance:** Continuously monitor query performance and make adjustments as needed.
- **Identify Slow Queries:** Use tools to identify and analyze slow-running queries for further optimization.

By applying these strategies, you can significantly enhance the performance of your database queries during implementation.

what would happen if stage in database management is ignored?

Ignoring any stage in the **Database Life Cycle Management** can lead to significant issues. Here's what might happen if each stage is overlooked:

1. Ignoring Requirements Analysis:

- **Misalignment with Needs:** The database may not meet user needs, leading to dissatisfaction and underutilization.
- **Wasted Resources:** Time and resources may be spent on features that are not necessary or relevant.

2. Ignoring Design and Plan:

- **Poor Structure:** Without a proper design, the database may become disorganized, making it difficult to retrieve and manage data.
- **Scalability Issues:** The database may not be able to grow or adapt to future needs, leading to performance bottlenecks.

3. Ignoring Implementation:

- **Data Migration Problems:** Skipping this stage can result in data loss or corruption during migration.
- **Security Vulnerabilities:** Without proper implementation, security measures may be inadequate, exposing the database to threats.

4. Ignoring Monitor and Maintain:

- **Performance Degradation:** Over time, the database may slow down due to unoptimized queries or lack of maintenance.
- **Increased Downtime:** Operational issues may arise, leading to system failures and downtime, which can affect business operations.
- **Data Breaches:** Without regular monitoring, security vulnerabilities may go unnoticed, increasing the risk of data breaches.

In summary, neglecting any stage can compromise the effectiveness, security, and reliability of the database, ultimately impacting the overall success of the project.

1. Among the first things that database administrators do each day is to check the databases they own for overall status. Which of the following is the best tool for this task?

1 / 1 point

☒ Dashboards

Correct. DBAs use preconfigured dashboards to make sure the most crucial information is quickly and easily available.

- ☐ Individual database server logs
- ☐ Key performance indicators
- ☐ Backup jobs

2. Which is an essential part of the Design and Plan stage of the database life cycle?

☐ Monitoring and maintenance

☒ Database modelling

Correct. In working on a database plan, it's important to model the database to assess how the parts will interact, how users will access the database, and so on.

- ☐ Automating repeating tasks
- ☐ Identifying users

3. Database administrators provide user support for the databases they're responsible for. Which of the following methods are best used to notify DBAs of typical or common user problems?

- ☐ Phone calls
- ☐ Texting
- ☐ Email
- ☒ Support tickets

Correct. Support tickets from self-service portals or a help desk are issued for DBA support because they track actions taken and other history of the issue.

3. The data life cycle consists of four stages. In which of these stages do stakeholders clarify the goals of the database?

- ☐ Design and plan
- ☒ Requirements analysis

Correct. This is the stage that includes clarifying the goals the database will fulfill.

- ☐ Monitoring and Maintenance
- ☐ Implementation

5. What is a best practice for compliance?

- ☐ Be sure that none of your users can access the database without the proper permissions.
- ☒ Check to make sure you understand all the regulations that may affect your organization on the national, international, and industry levels.

Correct. Failure to comply with all relevant laws and specifications may result in loss of business, data insecurity, or even legal action.

- ☐ Make sure you have the credentials and certifications you need to work as a DBA.
- ☐ Write a report every day about what you have done and the time you spent on it.

1. When a DBA meets with stakeholders like developers, data engineers, and architects to discuss database design, which of the following might be discussed?

- ☐ What platform will be used to develop the database software
- ☐ The cost of server resources
- ☐ Which logs will be reviewed daily
- ☒ Stress test scenarios for measuring the database's maximum load

✓ **Correct**

Correct, DBAs plan how to test the database's load by planning how to stress it.

2. Which is an essential part of the Design and Plan stage of the database life cycle?

- ☒ Database modeling
- ☐ Identifying users
- ☐ Automating repeating tasks
- ☐ Monitoring and maintenance

✓ **Correct**

Correct. In working on a database plan, it's important to model the database to assess how the parts will interact, how users will access the database, and so on.

Question 1

When a DBA meets with stakeholders like developers, data engineers, and architects to discuss database design, which of the following might be discussed?

✓ **Stress test scenarios for measuring the database's maximum load**

(✓ This is a key topic in database design discussions, ensuring the database can handle expected workloads.)

✗ Other options:

- **What platform will be used to develop the database software** → (Not typically a DBA's concern; more of a development decision.)
- **The cost of server resources** → (More relevant to infrastructure planning, not necessarily design discussions.)
- **Which logs will be reviewed daily** → (Operational concern, not a design discussion topic.)

Question 2

Which is an essential part of the Design and Plan stage of the database life cycle?

✓ **Database modeling**

(✓ This is the core activity where schemas, relationships, and structures are defined.)

✗ Other options:

- **Identifying users** → (Done in Requirements Analysis, not Design & Plan.)
- **Automating repeating tasks** → (Part of Monitoring & Maintenance, not Design.)
- **Monitoring and maintenance** → (Happens post-implementation, not during design.)



4. What is transparency?

- ☐ Letting all users know who is using the system at all times.
- ☐ Using color coding to make the user interface easier to understand.
- ☐ Regular maintenance and backup of the database.
- ☒ Making sure owners of information know what data you plan to collect and how you will use it.

✓ **Correct**

Correct. Transparency is an important ethical practice that builds trust and respect.

5. What is a best practice for compliance?

- ☐ Always host your database in the cloud.
- ☐ Write a report every day about what you have done and the time you spent on it.
- ☐ Make sure you follow laws, but don't worry about industry standards.
- ☒ Check to make sure you understand all the regulations that may affect your organization on the national, international, and industry levels.

✓ **Correct**

Correct. Failure to comply with all relevant laws and specifications may result in loss of business, data insecurity, or even legal action.

Question 4

What is transparency?

✅ Making sure owners of information know what data you plan to collect and how you will use it.

(✅ Transparency means clear communication about data usage, per privacy laws like GDPR.)

❌ Other options:

- Letting all users know who is using the system at all times → (Privacy risk, not transparency.)
- Using color coding for UI → (Accessibility, not transparency.)
- Regular maintenance and backup → (Operational task, unrelated to transparency.)

Question 5

What is a best practice for compliance?

✅ Check to make sure you understand all the regulations that may affect your organization on the national, international, and industry levels.

(✅ Compliance requires awareness of all applicable laws and standards.)

❌ Other options:

- Always host your database in the cloud → (Compliance depends on controls, not hosting location.)
- Write a daily report → (Excessive; compliance is about adherence, not documentation frequency.)
- Follow laws but ignore industry standards → (Standards often enforce best practices beyond law)

What is the role of schemas in database management?

Schemas play several important roles in database management:

1. Organization:

- Schemas provide a logical structure for organizing database objects, such as tables, views, and indexes. This organization helps users and administrators easily locate and manage these objects.

2. Naming Context:

- Schemas define a naming context for database objects, preventing naming conflicts. For example, two tables can have the same name if they belong to different schemas, allowing for better organization.

3. Access Control:

- Schemas help manage security by allowing administrators to set permissions at the schema level. This means specific users can be granted access to certain schemas while restricting access to others, enhancing data security.

4. Data Integrity:

- By grouping related objects within a schema, it becomes easier to enforce data integrity rules and constraints, ensuring that the data remains accurate and consistent.

5. Simplified Management:

- Schemas simplify database management tasks, such as backups, migrations, and updates. Administrators can manage objects within a schema collectively, rather than handling each object individually.

6. Support for Multiple Applications:

- In environments where multiple applications access the same database, schemas allow for separation of objects used by different applications, reducing the risk of interference and confusion.

In summary, schemas are essential for organizing, securing, and managing database objects effectively, contributing to overall database performance and integrity.

What is the difference between physical and logical storage?

The difference between **physical storage** and **logical storage** in databases can be understood as follows:

- **Physical Storage:**

- Refers to how data is actually stored on the hardware, such as hard drives or SSDs.
- It involves the specific files, directories, and devices that hold the data.
- For example, a physical storage might be a specific file on a disk that contains the data for a table.

- **Logical Storage:**

- Refers to how data is organized and accessed from the user's perspective, without needing to know where it is physically stored.
- It involves the structure and relationships of the data, such as tables, indexes, and views.
- For example, when you query a database to retrieve data from a table, you are interacting with the logical structure, not worrying about the underlying physical files.

In summary, physical storage is about the actual location of data on hardware, while logical storage is about how that data is organized and accessed in a way that makes sense to users and applications.

How can logical storage improve database management?

Logical storage can improve database management in several ways:

- **Abstraction:** It allows database administrators (DBAs) to manage data without needing to know the details of physical storage. This abstraction simplifies tasks like backups and data retrieval.
- **Flexibility:** Logical storage enables changes to the database structure (like adding or modifying tables) without affecting the physical storage. This means

you can reorganize data for better performance without worrying about where it is stored.

- **Easier Maintenance:** With logical storage, DBAs can perform operations like backups and restores using logical objects (like tables) rather than dealing with physical files. This makes maintenance tasks more straightforward and less error-prone.
- **Performance Optimization:** Logical storage allows for better organization of data. For example, frequently accessed data can be logically grouped, making it easier to optimize performance by placing it on faster storage.
- **Data Integrity:** Logical storage helps maintain relationships between data objects, ensuring that data integrity is preserved even when the underlying physical storage changes.

By leveraging logical storage, DBAs can manage databases more efficiently, leading to improved performance and easier maintenance.

1. Relational Database Management Systems (RDBMS) are organized in a hierarchical manner. What is usually the first object in the hierarchy?

☒ Instance

Correct. Most RDBMS begin with an instance, a single way of organizing the database and all it contains.

☐ Schemas

☐ Indexes

☐ Tables

2. RDBMSs store database metadata in special objects. Which of these is an example of metadata?

- ☐ Index
- ☐ Table
- ☒ Access privileges

Correct. Access privileges are metadata, as are the name of a table or the data type of a column.

- ☐ Db2

5. There are several distinct benefits of using tablespaces. Which of the following benefits removes the need to know the storage container that stores each object or tablespace?

- ☐ Performance
- ☒ Recoverability

Correct. A single command can back up or restore all database objects without worrying about which storage container each object of the tablespace is stored on.

- ☐ Authorization
- ☐ Storage management

4. Which of these is a way to expand storage space in cloud-based databases?

- ☐ Different disks ✓
- ☐ APIs
- ☒ Logical objects

Incorrect. A logical object is a way to refer to data no matter where it resides physically.

- ☐ External drives ✓

1. Which of the following schemas is the default schema?

- ☐ Object schema
- ☒ User schema
- ☐ System schema
- ☐ Partition schema

✓ **Correct**

Correct. The default schema is the user schema for the user who is currently logged on.

2. Which of the following is one of the four MYSQL system databases?

- ☐ Database 1
- ☐ Catalog
- ☒ Sys
- ☐ Read-only table

✓ **Correct**

Correct. Sys is one of the four system databases in each new MySQL instance.

Question 1

Which of the following schemas is the default schema?

☒ User schema

(✓ In most RDBMS, the default schema is typically the user schema where objects are created unless specified otherwise.)

☐ Other options:

- Object schema → (Not a standard default schema.)
- System schema → (Contains system objects, not user-created ones.)
- Partition schema → (Used for partitioning, not default storage.)

Question 2

Which of the following is one of the four MySQL system databases?

☒ Sys

(✓ The `sys` database in MySQL is a system database used for diagnostics and performance tuning.)

☐ Other options:

- Database 1 → (Not a standard MySQL system database.)
- Catalog → (Not a MySQL system database.)
- Read-only table → (Not a database; describes table permissions.)

Note: The four main MySQL system databases are:

1. `mysql` (stores user privileges)
2. `information_schema` (metadata)
3. `performance_schema` (performance metrics)
4. `sys` (simplifies diagnostics).



3. Some RDBMSs provide storage groups that let you organize data and storage based on data groups called Hot data, Warm data, and Cold data. Which temperature contains data that is accessed least frequently?

1 /

- ☐ Hot data
- ☒ Cold data
- ☐ Variable data
- ☐ Warm data

✓ **Correct**

Correct. Cold data is accessed less than Hot or Warm data.

4. Which of the following describes the relationship between containers and tablespaces?

1 /

- ☐ Tablespaces hold empty space inside containers
- ☐ Tablespaces store containers
- ☒ One or more containers store each tablespace
- ☐ Containers store multiple tablespaces

✓ **Correct**

Correct. Tablespaces are structures containing database objects and one or more containers store each tablespace.

Question 3

Which temperature contains data accessed least frequently?

☒ Cold data

(✓ Cold data is rarely accessed and often stored in cheaper, slower storage.)

✗ Other options:

- Hot data → (Frequently accessed, high-performance storage.)
- Warm data → (Moderately accessed, balanced storage.)
- Variable data → (Not a standard data temperature classification.)

Question 4

What describes the relationship between containers and tablespaces?

☒ One or more containers store each tablespace

(✓ Tablespaces are logical storage units that map to physical containers (files/directories).)

✗ Other options:

- Tablespaces hold empty space inside containers → (Incorrect; containers hold tablespace data.)
- Tablespaces store containers → (Backwards relationship.)
- Containers store multiple tablespaces → (A tablespace can span containers, not vice versa.)



Database partitions store subsets of data from very large databases. Which of the following is true?

- ☐ Tables cannot be partitioned into multiple logical partitions.
- ☒ A relational database whose data is managed across multiple partitions is called a partitioned relational database.
- ☐ Partitioning is used when little data is available.
- ☐ Each partition contains a copy of the entire set of data.

☒ Correct

Correct. "Partitioned relational database" is the name for this type of database.

Question 5

Which statement about database partitions is true?

✅ A relational database whose data is managed across multiple partitions is called a partitioned relational database.

(✅ Partitioning splits large tables into subsets for performance and manageability.)

❌ Other options:

- Tables cannot be partitioned into multiple logical partitions → (False; partitioning does this.)
- Partitioning is used when little data is available → (False; it's for large datasets.)
- Each partition contains a copy of the entire set of data → (False; partitions hold subsets.)

Summary of Correct Answers:

1. User schema
2. Sys
3. Cold data
4. One or more containers store each tablespace
5. A partitioned relational database manages data across partitions

What would happen if you couldn't restore a backup successfully?

If you couldn't restore a backup successfully, it could lead to several serious issues:

- **Data Loss:** You might lose important information that was stored in the database. This could affect business operations, customer data, or any critical information.
- **Downtime:** The system may be unavailable for a period while you try to resolve the issue, leading to interruptions in service and potentially lost

revenue.

- **Increased Recovery Time:** Without a successful backup, you may need to spend more time trying to recover data from other sources, which can be time-consuming and complicated.
- **Reputation Damage:** If customers or users are affected by the data loss or downtime, it could harm the reputation of the organization.
- **Compliance Issues:** For businesses that need to adhere to regulations regarding data protection, failing to restore backups could lead to legal consequences.

To avoid these issues, it's essential to regularly test your backup and restore processes to ensure they work as expected.

What would you do if a backup fails during restoration?

If a backup fails during restoration, here are some steps you can take to address the issue:

1. **Check Error Messages:** Review any error messages or logs generated during the restoration process. This can provide clues about what went wrong.
2. **Verify Backup Integrity:** Ensure that the backup file is not corrupted. You can do this by checking the file size and running integrity checks if available.
3. **Use an Alternative Backup:** If you have multiple backups, try restoring from a different backup file that was created at an earlier time.
4. **Consult Documentation:** Refer to the documentation for your database management system (DBMS) for troubleshooting tips specific to the restoration process.
5. **Seek Help:** If you're unable to resolve the issue, consider reaching out to a colleague or a database administrator for assistance.

6. **Test Your Backups Regularly:** After resolving the issue, it's a good practice to regularly test your backup and restore processes to ensure they work correctly in the future.

Taking these steps can help you recover from a failed restoration attempt and minimize the impact on your data and operations.

What is the difference between logical and physical backups?

The difference between **logical** and **physical backups** lies in how they capture and store data from a database:

Logical Backups:

- **Definition:** A logical backup extracts data and database structure in a format that can be used to recreate the database.
- **Content:** It includes Data Definition Language (DDL) commands (like `CREATE TABLE`) and Data Manipulation Language (DML) commands (like `INSERT`).
- **Granularity:** You can back up specific objects, such as individual tables or databases.
- **Restoration:** Restoring a logical backup typically creates a clean version of the database, reclaiming any wasted space.
- **Use Cases:** Useful for transferring data between different systems or for backing up specific parts of a database.

Physical Backups:

- **Definition:** A physical backup creates a copy of the actual storage files and directories that make up the database.
- **Content:** It includes all physical files, such as data files, configuration files, and log files.
- **Granularity:** You cannot easily restore individual tables; the backup is tied to the entire database structure.

- **Restoration:** Restoring a physical backup is generally faster and is often used for large databases requiring quick recovery.
- **Use Cases:** Ideal for complete system recovery or point-in-time recovery scenarios.

In summary, logical backups focus on the data and structure in a way that allows for selective restoration, while physical backups focus on the actual files and are better for complete system recovery.

What are the advantages of using incremental backups over differential backups?

Here are the advantages of using **incremental backups** over **differential backups**:

1. **Faster Backup Time:**

- Incremental backups only save the changes made since the last backup (whether full or incremental). This means they are generally quicker to perform compared to differential backups, which save all changes since the last full backup.

2. **Less Storage Space:**

- Because incremental backups only capture the changes, they typically require less storage space than differential backups, which can grow larger over time as they accumulate changes since the last full backup.

3. **Reduced Bandwidth Usage:**

- Incremental backups use less bandwidth when transferring data, making them more efficient for environments with limited network resources.

4. **Frequent Backups:**

- The speed and efficiency of incremental backups allow for more frequent backups throughout the day, which can lead to better data protection and less data loss in case of a failure.

However, it's important to note that while incremental backups have these advantages, restoring data from them can be more complex and time-consuming, as you need to restore the full backup first and then each incremental backup in order.

What challenges might arise when restoring from incremental backups?

Restoring from **incremental backups** can present several challenges:

1. **Complex Restoration Process:**

- You must first restore the last full backup and then apply each incremental backup in the correct order. If you miss any incremental backup, the restoration may fail or result in incomplete data.

2. **Time-Consuming:**

- The process of restoring multiple incremental backups can take a significant amount of time, especially if there are many incremental backups to apply.

3. **Dependency on Previous Backups:**

- If any incremental backup is corrupted or missing, it can hinder the entire restoration process. This means that the integrity of each incremental backup is crucial for a successful restore.

4. **Increased Risk of Errors:**

- The more steps involved in the restoration process, the higher the chance of human error. For example, accidentally applying backups in the wrong order can lead to data inconsistencies.

5. **Resource Intensive:**

- Restoring from multiple incremental backups may require more system resources (CPU, memory, and disk I/O), which can impact the performance of the system during the restoration process.

Understanding these challenges can help you plan better for backup and restoration strategies.

How can you determine the best backup policy for your needs?

To determine the best backup policy for your needs, consider the following factors:

1. **Data Importance:**

- Assess how critical your data is to your operations. If losing data would significantly impact your business, you may need more frequent backups.

2. **Data Change Frequency:**

- Identify how often your data changes. For example, customer orders may change frequently, requiring more regular backups compared to static data like product information.

3. **Availability Requirements:**

- Consider how important it is for users to access data without interruption. If your system needs to be available 24/7, hot backups might be necessary, while cold backups could be scheduled during off-peak hours.

4. **Recovery Time Objective (RTO):**

- Determine how quickly you need to restore data after a failure. If you need rapid recovery, hot backups may be more suitable.

5. **Recovery Point Objective (RPO):**

- Decide how much data loss is acceptable. If you can tolerate losing a few hours of data, less frequent backups may suffice.

6. **Backup Storage:**

- Evaluate where you will store backups. Hot backups are often stored on available servers, while cold backups may be on external drives or offline servers.

7. Automation Options:

- Check if your database management system (RDBMS) allows for automated backups. Automation can help ensure backups are performed regularly without manual intervention.

By analyzing these factors, you can create a backup policy that aligns with your operational needs and data management strategy.

How would you implement a backup policy for a 24/7 system?

Implementing a backup policy for a 24/7 system requires careful planning to ensure data availability and integrity without disrupting user access. Here are key steps to consider:

1. Use Hot Backups:

- Since the system operates continuously, opt for **hot backups** that allow data to be backed up while the database is active. This minimizes downtime and keeps services available.

2. Schedule Incremental Backups:

- Implement **incremental backups** frequently (e.g., hourly or every few hours). This captures only the changes made since the last backup, reducing the load on the system.

3. Perform Full Backups During Off-Peak Hours:

- Schedule **full backups** during off-peak hours (e.g., late at night or early morning) when user activity is minimal. This helps reduce the impact on performance.

4. Automate Backup Processes:

- Utilize automation features in your RDBMS to schedule and manage backups. This ensures consistency and reduces the risk of human error.

5. Monitor Backup Performance:

- Regularly monitor the performance of backup processes to ensure they do not degrade system performance. Adjust schedules or methods as needed.

6. Test Restore Procedures:

- Regularly test your restore procedures to ensure that backups can be successfully restored. This helps identify any issues before a real recovery is needed.

7. Consider Redundancy:

- Store backups in multiple locations (e.g., on-site and off-site or in the cloud) to protect against data loss from hardware failures or disasters.

8. Document the Backup Policy:

- Clearly document the backup policy, including schedules, procedures, and responsibilities. This ensures everyone involved understands their roles.

By following these steps, you can create a robust backup policy that supports a 24/7 system while ensuring data safety and availability.

How can you automate backup processes effectively?

To automate backup processes effectively, consider the following strategies:

1. Utilize Built-in RDBMS Features:

- Most Relational Database Management Systems (RDBMS) offer built-in scheduling tools. Use these features to set up regular backup schedules without manual intervention.

2. Scripted Backups:

- Write scripts (using languages like SQL, PowerShell, or Bash) to automate backup tasks. These scripts can be scheduled to run at specific times using task schedulers (e.g., cron jobs on Linux or Task Scheduler on Windows).

3. Use Backup Management Tools:

- Consider using third-party backup management tools that provide automation features. These tools often come with user-friendly interfaces and additional functionalities like monitoring and reporting.

4. Set Up Notifications:

- Configure notifications to alert you of backup successes or failures. This ensures you are informed of any issues that may require attention.

5. Implement Versioning:

- Automate the retention of multiple backup versions. This allows you to restore from different points in time, which can be crucial in case of data corruption or accidental deletions.

6. Test Automation Regularly:

- Periodically test the automated backup processes to ensure they are functioning correctly. This includes verifying that backups are being created and can be restored successfully.

7. Document Automation Procedures:

- Keep clear documentation of your automated backup processes, including scripts, schedules, and configurations. This helps in troubleshooting and onboarding new team members.

By implementing these strategies, you can create a reliable and efficient automated backup process that minimizes manual effort and reduces the risk of data loss.

How would you implement a backup schedule for a business?

Implementing a backup schedule for a business involves several key steps to ensure data integrity and availability. Here's a structured approach:

1. **Assess Business Needs:**

- Identify critical data and applications that require regular backups.
Understand the impact of data loss on business operations.

2. **Define Backup Frequency:**

- Determine how often backups should occur based on data importance and change frequency:
 - **Critical Data:** Consider hourly or daily backups.
 - **Less Critical Data:** Weekly or bi-weekly backups may suffice.

3. **Choose Backup Types:**

- Decide on the types of backups to implement:
 - **Full Backups:** Capture all data; schedule these less frequently (e.g., weekly).
 - **Incremental Backups:** Capture only changes since the last backup; schedule these more frequently (e.g., daily or multiple times a day).
 - **Differential Backups:** Capture changes since the last full backup; these can be scheduled between full backups.

4. **Schedule Backups:**

- Use off-peak hours for full backups to minimize impact on system performance. For example, schedule full backups during weekends or late at night.
- Schedule incremental backups during business hours if using hot backups, ensuring they do not disrupt user activities.

5. **Automate the Process:**

- Utilize RDBMS features or backup management tools to automate the backup process. This reduces the risk of human error and ensures consistency.

6. Monitor and Adjust:

- Regularly monitor backup performance and success rates. Adjust the schedule or methods as needed based on system performance and business requirements.

7. Test Restore Procedures:

- Periodically test the restore process to ensure backups can be successfully restored. This helps identify any issues before a real recovery is needed.

8. Document the Backup Schedule:

- Create clear documentation outlining the backup schedule, types, and responsibilities. This ensures everyone involved understands the process.

By following these steps, you can establish a comprehensive backup schedule that meets your business's needs and protects critical data.

1. Which of the following is not a scenario that might require backup and restore?

- ☐ Transfer data from one database to another
- ☐ Disaster recovery
- ☐ Facilitate a change of RDBMS
- ☒ When a new user logs in

Correct. It's not necessary to back up a database whenever a user logs in.

2. Which of the following is a benefit of logical backups?

- ☐ Can only restore to a similar RDBMS only
- ☒ Reclaims wasted space

Correct. Logical backup/restore reclaims any wasted space from the original database as restore creates a clean version of the tables.

- ☐ Copies configuration and log files
- ☐ Smaller and quicker than physical backups

3. Full backups are complete copies of all data in the objects you're backing up. When the database size increases, which of the following does not affect full backups?

☒ Latency

Correct. Latency doesn't affect full backups, rather the system network might.

- ☐ Time
- ☐ Storage
- ☐ Bandwidth

4. Which of the following is a drawback with using hot backups?

- ☐ Availability
- ☒ Performance degradation

Correct. Hot backups can cause performance degradation for users while backups are running.

- ☐ Users are logged out of the system
- ☐ Stored on external drives

5. Which of the following is not a consideration for setting a backup policy?

☒ Type of hardware

Correct. This may be a feature of designing the system, but it shouldn't affect backups.

- ☐ Data usage patterns
- ☐ Encryption
- ☐ Physical or logical backups

1. Which of the following explains one reason that you should schedule or automate your backups?

- ☐ Save storage space
- ☐ Stop data loss
- ☒ Faster backup
- ☐ Less user impact ✓

✗ **Incorrect**

Incorrect. Please review the Backup Policies video.

2. Which of these is a characteristic of hot backups?

- ☐ Cannot be used in 24/7 environments
- ☒ Has no impact on availability
- ☐ Stored on external drives
- ☐ Does not result in performance degradation for users

✓ **Correct**

Correct. Hot backups allow users to continue with their activities.

2. Characteristic of hot backups:

✓ **Answer: Has no impact on availability**

Hot backups occur while the database is running, allowing continuous availability (although there may still be minor performance impacts).

3. Full database backups restore data to the state it was in when the backup occurred. A way around this is to run a point-in-time backup/recovery. What is the first and most important step needed in point-in-time backups?

- ☒ Enable transaction logging
- ☐ Apply new full backup recovery
- ☐ Run DML statements before backups
- ☐ Store multiple copies of full backup

✔ **Correct**

Correct. Enabling the transaction logs on a database so you can use the information in the log file to reapply the transactions to restore the database.

4. Which of these is true about differential and incremental backups?

- ☐ They are essentially identical.
- ☒ Differential backs up data since the last full backup, while incremental backs up data since the last backup of any type.
- ☐ Incremental backups copy anything specified
- ☐ Incremental backups back up user data, while differential backups back up tables.

✔ **Correct**

Correct. this is how differential backups perform.

3. First step in point-in-time backup/recovery:

✔ **Answer: Enable transaction logging**

Without transaction logs, you can't roll forward to a specific point in time during recovery.

4. Difference between differential and incremental backups:

✔ **Answer: Differential backs up data since the last full backup, while incremental backs up data since the last backup of any type.**

This is the standard definition and key distinction between the two.

5. Which of these are important key considerations when using backup and restore?

- ☐ Correct cloud provider
- ☒ Validated backup files and functional restore plan
- ☐ Backup scheduling
- ☐ Storage level snapshots

✔ **Correct**

Correct. always check that the backup file or files are valid and the restore plan is successful.

5. Key considerations for backup and restore:

✔ **Answer: Validated backup files and functional restore plan**

Backups are only useful if they are valid and restorable. Testing your restore process is crucial.

What would happen if users had too many privileges?

If users have too many privileges in a database, several issues can arise:

1. **Security Risks:** Users with excessive privileges can accidentally or intentionally access, modify, or delete sensitive data. This can lead to data breaches or loss of critical information.
2. **Data Integrity Issues:** When users can make changes without restrictions, it increases the risk of errors. For example, a user might accidentally delete important records or alter data incorrectly, leading to inconsistencies.
3. **Compliance Violations:** Many industries have regulations regarding data access and protection. If users have more privileges than necessary, it could result in non-compliance with these regulations, leading to legal consequences.

4. **Increased Attack Surface:** If a user's account is compromised, an attacker could exploit their excessive privileges to cause significant damage, such as stealing data or disrupting services.

To mitigate these risks, it's essential to follow the **principle of least privilege**, which means granting users only the access they need to perform their job functions.

How can you implement least privilege in a database?

Implementing the **principle of least privilege** in a database involves several key steps:

1. **Role-Based Access Control (RBAC):**

- Create roles based on job functions and assign permissions to these roles rather than to individual users. This way, users only receive the permissions necessary for their specific role.

2. **Regularly Review Permissions:**

- Periodically audit user permissions to ensure they still align with their current job responsibilities. Remove any unnecessary privileges.

3. **Limit Administrative Access:**

- Restrict administrative privileges to a small number of trusted users. Regular users should not have administrative rights unless absolutely necessary.

4. **Use Object-Level Permissions:**

- Grant permissions at a granular level, such as on specific tables or columns, rather than giving blanket access to entire databases.

5. **Implement Temporary Privileges:**

- For tasks that require elevated access, consider granting temporary permissions that expire after the task is completed.

6. **Monitor and Audit Access:**

- Continuously monitor user activity and maintain logs of access and changes. This helps identify any unauthorized access or misuse of privileges.

7. Educate Users:

- Train users on the importance of data security and the principle of least privilege, ensuring they understand their responsibilities regarding access.

By following these steps, you can effectively implement the principle of least privilege in your database, enhancing security and reducing risks.

How can you manage users, groups, and roles effectively?

To manage users, groups, and roles effectively in a database, you can follow these key practices:

1. Define Clear Roles:

- Create specific roles based on job functions (e.g., admin, backup operator, salesperson).
- Ensure each role has only the permissions necessary for its tasks (principle of least privilege).

2. Use Groups for Similar Users:

- Group users with similar access needs together (e.g., all sales staff).
- Assign permissions to the group rather than to each individual user, simplifying management.

3. Regularly Review Permissions:

- Periodically check user roles and group memberships to ensure they still align with current job functions.
- Remove access for users who no longer need it (e.g., when they change roles or leave the organization).

4. Implement External Authentication:

- Use external systems (like LDAP or Kerberos) for user authentication to centralize user management and enhance security.

5. Document User Management Policies:

- Keep clear documentation of user roles, groups, and permissions.
- This helps in onboarding new users and maintaining consistency in access management.

6. Automate Where Possible:

- Use scripts or tools to automate the assignment of roles and permissions, especially for large organizations.

By following these practices, you can ensure that user access is managed efficiently and securely.

How can you automate user role assignments in a database?

Automating user role assignments in a database can streamline management and enhance security. Here are some common methods to achieve this:

1. Scripting:

- Use SQL scripts to automate the creation and assignment of roles. For example, you can write a script that assigns a specific role to all users in a certain group.

```
GRANT role_name TO user_name;
```

2. Database Management Tools:

- Many database management systems (DBMS) offer built-in tools or features for role management. You can schedule tasks to run at specific intervals to update user roles based on predefined criteria.

3. External Identity Management Systems:

- Integrate with identity management solutions (like LDAP or Active Directory) that can automatically assign roles based on user attributes (e.g., department, job title).

4. Triggers:

- Implement database triggers that automatically assign roles when a new user is created or when a user's attributes change. This ensures that role assignments are always up-to-date.

5. Scheduled Jobs:

- Use database job schedulers to run scripts at regular intervals. For example, a job could check for new users and assign them roles based on their department.

6. APIs:

- If your database supports it, use APIs to programmatically manage user roles. This allows for integration with other systems and can automate role assignments based on external events.

By implementing these methods, you can ensure that user role assignments are consistent and efficient.

What would happen if roles have excessive permissions assigned?

If roles have excessive permissions assigned, several issues can arise:

1. Security Risks:

- Users may gain access to sensitive data or critical functions that they do not need for their job. This increases the risk of data breaches or unauthorized actions.

2. Increased Attack Surface:

- If a user's account is compromised, the attacker could exploit the excessive permissions to cause significant damage, such as deleting data or altering critical configurations.

3. Compliance Issues:

- Many industries have regulations regarding data access and security. Excessive permissions can lead to non-compliance, resulting in legal penalties or fines.

4. Operational Errors:

- Users with too many permissions might accidentally perform actions that disrupt operations, such as modifying or deleting important database objects.

5. Difficulty in Management:

- Managing roles with excessive permissions can become complex and confusing, making it harder to track who has access to what and complicating audits.

6. Principle of Least Privilege Violation:

- Assigning excessive permissions goes against the principle of least privilege, which states that users should only have the minimum access necessary to perform their job functions. This can lead to a culture of poor security practices.

To mitigate these risks, it's essential to regularly review and adjust role permissions, ensuring they align with users' current job responsibilities.

How can you assess role permissions effectively?

To assess role permissions effectively, you can follow these steps:

1. Inventory of Roles and Permissions:

- Create a comprehensive list of all roles and their associated permissions within the database. This serves as a baseline for assessment.

2. Define Job Functions:

- Clearly outline the responsibilities and tasks associated with each job function in your organization. This helps in determining the necessary permissions for each role.

3. Conduct Regular Audits:

- Schedule periodic audits to review role permissions. Check if the permissions assigned align with the current job functions and responsibilities of users.

4. Utilize Access Logs:

- Analyze access logs to monitor how users interact with the database. Identify any unusual access patterns or permissions that are not being used.

5. Engage Stakeholders:

- Involve department heads or team leaders in the assessment process. They can provide insights into the actual needs of their teams and help identify any discrepancies.

6. Implement the Principle of Least Privilege:

- Ensure that roles are designed to provide only the minimum permissions necessary for users to perform their tasks. Remove any excessive permissions identified during the assessment.

7. Document Changes:

- Keep detailed records of any changes made to role permissions during the assessment process. This documentation can be useful for future audits and compliance checks.

8. Feedback Mechanism:

- Establish a process for users to provide feedback on their access needs. This can help identify any gaps or unnecessary permissions.

By following these steps, you can ensure that role permissions are appropriate and aligned with organizational needs.

What is the purpose of granting permissions in databases?

The purpose of granting permissions in databases is to **control access** to data and ensure **security**. Here are the key reasons:

- **Data Security:** By granting permissions, you can protect sensitive information from unauthorized access. Only users with the right permissions can view or modify data.
- **Role Management:** Permissions can be assigned to users, groups, or roles, allowing for easier management of who can do what within the database. This helps in organizing access based on job functions.
- **Data Integrity:** By controlling who can insert, update, or delete data, you help maintain the accuracy and consistency of the data in the database.
- **Audit and Compliance:** Granting permissions allows for tracking who accessed or modified data, which is important for compliance with regulations and auditing purposes.
- **Operational Efficiency:** By granting specific permissions, you can streamline workflows, allowing users to perform their tasks without unnecessary barriers.

What is the difference between granting and revoking permissions?

The difference between granting and revoking permissions in databases is as follows:

- **Granting Permissions:**
 - **Definition:** This is the process of giving a user, group, or role the ability to perform specific actions on database objects (like tables).
 - **Purpose:** It allows users to access and manipulate data according to their needs. For example, granting "select" permission lets a user read data from a table.

- **Example:** Using the SQL command `GRANT SELECT ON mytable TO user1;` gives user1 permission to read data from "mytable."
- **Revoking Permissions:**
 - **Definition:** This is the process of taking away previously granted permissions from a user, group, or role.
 - **Purpose:** It restricts access to data when a user no longer needs it or when their role changes. This helps maintain security and data integrity.
 - **Example:** Using the SQL command `REVOKE SELECT ON mytable FROM user1;` removes user1's ability to read data from "mytable."

In summary, granting permissions allows access, while revoking permissions removes access.

What happens if you deny a permission after granting it?

When you **deny a permission** after it has been granted, the denial takes precedence over the grant. Here's what happens:

- **Override of Granted Permissions:** The deny command will override any previously granted permissions for that specific action. This means that even if a user was granted access to perform an action (like reading data), the deny command will block that access.
- **Immediate Effect:** The effect of denying a permission is immediate. Once the permission is denied, the user will no longer be able to perform the action, regardless of any other permissions they may have.
- **Example:** If a user is granted "SELECT" permission on a table and then a "DENY SELECT" command is issued, the user will not be able to read data from that table anymore, even though they were previously granted access.

This mechanism is useful for ensuring that certain users cannot access specific data, regardless of their other permissions.

How can you effectively manage user permissions in a database?

Effectively managing user permissions in a database involves several best practices:

1. **Role-Based Access Control (RBAC):**

- Create roles that group permissions based on job functions (e.g., admin, user, guest).
- Assign users to these roles instead of granting permissions individually, simplifying management.

2. **Principle of Least Privilege:**

- Grant users only the permissions they need to perform their tasks. This minimizes the risk of unauthorized access or accidental data changes.

3. **Regular Audits and Reviews:**

- Periodically review user permissions to ensure they are still appropriate. Remove or adjust permissions for users who no longer need access.

4. **Use of Groups:**

- Organize users into groups based on their roles or departments. Grant permissions to groups rather than individuals to streamline management.

5. **Documentation:**

- Maintain clear documentation of permissions granted, roles created, and any changes made. This helps in tracking and auditing access.

6. **Testing Changes:**

- Before implementing significant changes to permissions, test them in a controlled environment to ensure they work as intended without disrupting access for legitimate users.

7. **Implementing Deny Permissions:**

- Use the deny command judiciously to explicitly block access to sensitive data for certain users or roles, overriding any granted permissions.

By following these practices, you can maintain a secure and organized approach to managing user permissions in a database.

What are the trade-offs between encryption and performance?

The trade-offs between **encryption** and **performance** can be summarized as follows:

- **Performance Impact:**
 - **Encryption processes** can slow down system performance because they require additional computational resources. Encrypting and decrypting data takes time, which can lead to delays in data access and processing.
- **Algorithm Complexity:**
 - **Asymmetric encryption** (like RSA) generally has a greater performance impact than **symmetric encryption** (like AES) because it uses longer keys and more complex mathematical operations. This means that if you need faster performance, symmetric encryption is often preferred.
- **Data Size:**
 - The size of the data being encrypted can also affect performance. Larger datasets take longer to encrypt and decrypt, which can slow down applications that rely on quick data access.
- **System Resources:**
 - The available system resources (CPU, memory) can influence how much encryption affects performance. Systems with more powerful hardware can handle encryption tasks more efficiently.
- **Trade-off Decisions:**

- Organizations must balance the need for security with the need for performance. In some cases, they may choose to encrypt only sensitive data or use less complex encryption methods to maintain better performance.

In summary, while encryption enhances security, it can also lead to slower performance. Choosing the right type of encryption and understanding its impact on system performance is crucial for effective database management.

I

What is the impact of encryption on database performance?

The impact of **encryption on database performance** can be significant and can be summarized in the following points:

- **Increased Latency:**
 - Encrypting and decrypting data adds processing time, which can lead to increased latency when accessing or storing data. This means that queries may take longer to execute.
- **Resource Consumption:**
 - Encryption requires additional CPU and memory resources. This can strain system resources, especially in environments with high transaction volumes or limited hardware capabilities.
- **Complexity of Operations:**
 - Operations such as backups, restores, and data migrations may become more complex and time-consuming due to the need to encrypt or decrypt data during these processes.
- **Trade-offs with Security:**
 - While encryption enhances security, it may require organizations to make trade-offs regarding performance. For example, they might choose to

encrypt only sensitive data or use symmetric encryption for better performance.

- **Impact on Indexing:**

- Encrypted data may not be indexed as efficiently, which can affect query performance. This is particularly relevant for databases that rely heavily on indexing for fast data retrieval.

- **Configuration and Management:**

- Proper configuration and management of encryption settings are essential to minimize performance impacts. Some databases offer features like Transparent Data Encryption (TDE) to simplify encryption without significantly affecting performance.

In summary, while encryption is crucial for protecting sensitive data, it can lead to slower database performance and increased resource usage. Organizations need to carefully consider these factors when implementing encryption strategies.

1. What is authentication?

1 / 1 point

- ☒ A process of verifying that a user is who they say they are

Correct. Users must present a credential such as a password or a fingerprint to authenticate their access.

- ☐ Reviewing and using system-specific security features
- ☐ Maintaining security in a managed cloud environment
- ☐ Preventing insecure code.

2. Which of the following is true of groups and roles?

1 / 1 point

- ☐ When assigning groups and roles, add users to all the groups and roles they may ever need, since you can't change them later.
- ☒ Groups and roles simplify user management

Correct. Many functions can be handled more easily if users are assigned to groups and roles.

- ☐ Only pre-determined roles are available
- ☐ You must always define groups in the database to manage users

3. Privileges are given to users, groups, and roles. What defines a user's overall privileges?

1 / 1 point

- ☐ Roles permissions
- ☐ Groups permissions
- ☒ Combination of user-level permissions and those of their groups and roles

Correct. The combination of a user's personal permissions and those of all groups and roles they belong to defines their overall permissions.

- ☐ User's personal permissions

4. Which is a reason to audit a database?

1 / 1 point

- ☒ Identify gaps in security access and activity

Correct. Auditing can help identify gaps in your security system.

- ☐ Protect databases
- ☐ See if a database is cost-effective.
- ☐ Database integrity

5. At which point should you encrypt data?

1 / 1 point

- ☐ While data is archived
- ☒ At rest and transmission

Correct. Encrypting data can help protect data at rest and during transmission.

- ☐ When data is online
- ☐ When data is compromised

1. The OS for database servers must be secure. Which of the following methods that the video recommends, involves vulnerabilities?

- ☐ Access monitoring
- ☒ Hardened
- ☐ Regular patching
- ☐ Firewall implementation

✓ **Correct**

Correct. You should ensure your database server OS is hardened using a known configuration to reduce vulnerabilities.

2. When an RDBMS doesn't support user groups, how can you define groups?

- ☐ Map to user initials
- ☐ Additional user accounts
- ☒ Map to OS groups
- ☐ Alias accounts

✓ **Correct**

Correct. You can map a database group to an administrative group in the operating system.

Question 1

The OS for database servers must be secure. Which of the following methods that the video recommends, involves vulnerabilities?

✓ **Hardened**

(While hardening improves security, misconfigurations during the process can introduce vulnerabilities.)

Question 2

When an RDBMS doesn't support user groups, how can you define groups?

✓ **Map to OS groups**

(If the RDBMS lacks native group support, mapping to OS groups is a common workaround.)

3. Using SQL commands, how can you remove database permissions from a user, group, or role?

- ☐ Eliminate connect
- ☒ Revoke connect
- ☐ Delete connect
- ☐ Cancel connect

✓ **Correct**

Correct. The revoke connect to "user, group, or role" revokes permissions to databases.

4. Why should you audit failed attempts to access databases?

- ☒ Identify potential attacks
- ☐ Test database security
- ☐ Identify unauthorized users
- ☐ Add security levels

✓ **Correct**

Correct. Tracking failed attempts to access databases can help you identify potential attacks, like brute force attempts.

Question 3

Using SQL commands, how can you remove database permissions from a user, group, or role?

✓ **Revoke connect**

(The correct SQL command is **REVOKE** to remove permissions.)

Question 4

Why should you audit failed attempts to access databases?

✓ **Identify potential attacks**

(Failed access attempts may indicate brute-force attacks or unauthorized access attempts.)

5. Which of the following is true of asymmetric encryption?

- ☐ Key is shared with all users
- ☐ Is more susceptible to compromise than symmetric encryption
- ☒ Uses two keys: one public, one private
- ☐ DES is an example of asymmetric encryption

✔ **Correct**

Correct. The public key encrypts, while the private key decrypts.

Question 5

Which of the following is true of asymmetric encryption?

✔ **Uses two keys: one public, one private**

(Asymmetric encryption relies on a public key for encryption and a private key for decryption.)